

The Mechanics and Ethics Ethics of Large Language Language Models

Understanding the architecture, training, and societal implications of transformative AI systems

www.drshetty.in



What Makes a Language Model "Large"?

Large Language Models represent a sophisticated subset of Generative AI and Foundation Models. Whilst foundation models provide a versatile base for various adaptable outputs across multiple domains, LLMs focus specifically on processing and generating text and text-like data, including natural language and computer code.

The term "large" in LLMs encompasses three fundamental dimensions: the massive volume of training data measured in trillions of tokens, the extraordinary number of parameters that capture nuanced patterns, and the immense computational resources required for training. These dimensions work synergistically to create AI systems capable of remarkably human-like text generation.



Key Distinction

LLMs are foundation models specialised for language tasks, distinguishing them from models designed for images, audio, or multimodal applications.

The First Dimension: Unprecedented Data Volume

Training sets for modern LLMs consist of trillions of tokens—the fundamental units of text processing. This data is typically unlabelled and processed through self-supervised learning techniques, where the model learns patterns without explicit human annotation for every piece of information.



Common Crawl

A massive, continuously updated archive capturing vast swathes of the entire web, providing diverse linguistic patterns and knowledge domains.



Books

Comprehensive digital libraries containing published works, offering refined language structures and narrative coherence.



GitHub Repositories

Vast collections of computer code across programming languages, enabling models to understand technical syntax and logic.

Training datasets often reach petabytes in size. To contextualise: one gigabyte stores approximately 178 million words, whilst a petabyte represents one million gigabytes—an almost incomprehensible volume of textual information.

The Second Dimension: Parameter Count and Complexity

Parameters are the fundamental internal values—essentially weights and biases—that a model independently adjusts during the learning process. These numerical values determine how input data is transformed through the network's layers to produce outputs.

There exists a direct correlation between parameter count and model capability: more parameters enable the capture of increasingly nuanced patterns, subtle contextual relationships, and complex linguistic structures. However, this relationship is not purely linear—architectural innovations and training techniques also play crucial roles.

175B

GPT-3 Parameters

Processing a 45-terabyte training corpus



Helpful Analogy

If a model is a giant calculator, parameters are the billions of individual dials and settings that collectively determine the final computational result.

The Third Dimension: Massive Computational Requirements

Hardware Infrastructure

Training requires thousands of Graphics Processing Units (GPUs) working in parallel, coordinated through sophisticated distributed computing frameworks.

Temporal Investment

The training process spans months of continuous, high-intensity computation—a period during which the model gradually "learns" patterns from data.

Energy Consumption

The electricity demands are staggering, with implications for both operational costs and environmental sustainability that must be carefully considered.

These computational requirements represent a significant barrier to entry, concentrating LLM development capabilities amongst well-resourced organisations and institutions. The infrastructure needed to "bake" a state-of-the-art model rivals that of traditional supercomputing facilities.




The Transformer Revolution: A Paradigm Shift

The Pre-2017 Limitation

Before the Transformer architecture emerged in 2017, AI systems processed text sequentially—word by word, left to right. This created a fundamental limitation: as sentences grew longer, earlier context would gradually fade from the model's "memory," much like human short-term memory degradation.

Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks attempted to address this, but struggled with truly long-range dependencies.



The Transformer architecture fundamentally changed how AI understands language by enabling models to consider all words simultaneously, rather than processing them one at a time.

Self-Attention: The Brain of the LLM

The self-attention mechanism represents the revolutionary "brain" of modern LLMs. This component allows the model to dynamically weigh the importance of different words in a sentence relative to each specific word being processed—a capability that creates surprisingly sophisticated contextual understanding.

01	02	03
Relationship Computation	Importance Weighting	Contextual Representation
The model calculates attention scores between every word and every other word in the input sequence.	These scores determine how much "focus" each word should receive when understanding another word's meaning.	Words are transformed into context-aware representations that capture their meaning within the specific sentence.

By considering every word in relation to every other word, rather than processing sequentially, the model builds a comprehensive, holistic understanding of sentence structure, grammatical relationships, and semantic meaning. This parallel processing capability is what enables Transformers to capture long-range dependencies that earlier architectures missed.

Context Windows and Positional Encoding

The Context Window

Think of the context window as the model's "short-term memory"—it defines the maximum number of tokens the model can process and consider simultaneously in a single forward pass. Modern models have progressively expanded these windows from a few thousand tokens to hundreds of thousands.

This constraint creates a fundamental limitation: information beyond the context window becomes invisible to the model, regardless of its relevance. Techniques like retrieval-augmented generation (RAG) help circumvent this limitation.


Positional Encoding

Because Transformers process all words simultaneously through parallelisation, they inherently lack information about word order. To address this, mathematical "tags" called positional encodings are added to each token.

These encodings inject information about token position into the input, enabling the model to distinguish between "The cat chased the dog" and "The dog chased the cat"—sentences with identical words but different meanings due to word order.

Tokenisation: How Models "Read" Text

LLMs do not process individual characters or even complete words in the way humans read. Instead, they operate on tokens—variable-length chunks of characters that represent the fundamental units of processing. Understanding tokenisation is crucial to comprehending how these models actually interpret language.



Tokenisation Examples


- **"Hamburger"** might be treated as a single token, as it's a common word
- **"Unbelievable"** might split into ["un-", "believe-", "-able"], reflecting morphological structure
- **"AI"** could be one token, whilst **"Antidisestablishmentarianism"** would require multiple tokens

Why This Matters

Token boundaries affect model understanding and generation. Rare words split into more tokens, potentially degrading comprehension. This is why models sometimes struggle with very specialised terminology or neologisms.

The Three-Stage Learning Journey

The development of an LLM functions like a master chef's educational progression, advancing through three distinct and essential stages. Each stage builds upon the previous one, transforming a random initialisation into a sophisticated, helpful assistant capable of following complex instructions.



This three-stage approach represents decades of machine learning research distilled into a practical methodology. Each stage addresses different aspects of model capability, from raw linguistic competence to nuanced social appropriateness.

Stage 1: Pre-training—The Foundation

Pre-training represents the most computationally intensive phase, where the model learns general language structure, grammar, facts about the world, and reasoning patterns from massive unlabelled datasets. This stage creates what's called a "base model"—intellectually capable but often difficult to interact with directly.

The Core Mechanism: Next-Word Prediction

The training objective is deceptively simple yet remarkably powerful: given a sequence of words, predict what comes next. The model hides a word in a sentence and attempts to guess it. Initially, predictions are random ("The sky is bug"), but through countless iterations and error correction, the model learns statistical patterns that make predictions increasingly accurate ("The sky is blue").



Self-Supervised Learning

Pre-training requires no manual labelling. The data itself provides the supervision signal: the next word in each sequence serves as the training target, enabling learning from trillions of examples.

Pre-training is where the model develops its "intelligence"—the vast knowledge base and linguistic capabilities that underpin all later specialisation.

Stage 2: Supervised Fine-Tuning (SFT)

Whilst pre-training creates a knowledgeable base model, that model often produces outputs that are verbose, unfocused, or fail to follow instructions properly. Supervised Fine-Tuning bridges this gap, teaching the model to function as a helpful, instruction-following assistant.

The Training Methodology

Human annotators create thousands of high-quality prompt-response pairs demonstrating desired behaviour. For example: **Prompt:** "Explain photosynthesis simply" → **Response:** A clear, concise explanation appropriate for the intended audience.

Learning Instruction-Following

The model learns to recognise instruction formats and respond appropriately. It develops an understanding of different request types: questions requiring factual answers, tasks requiring step-by-step procedures, creative requests, and more.


Behavioural Alignment

Beyond technical capability, SFT installs desired behavioural characteristics: conciseness, clarity, appropriate tone, acknowledgment of uncertainty, and refusal of harmful requests.

This stage typically uses far less data than pre-training—tens of thousands rather than trillions of examples—but the quality of these examples is paramount. Poor SFT data can severely limit model usability despite excellent pre-training.

Stage 3: RLHF—Refining Through Feedback

Reinforcement Learning from Human Feedback (RLHF) represents the final refinement stage, addressing a fundamental challenge: human preferences are often difficult to specify explicitly but easy to demonstrate through comparative judgements.




RLHF is why modern LLMs feel remarkably human-like in their interactions—they've been extensively trained to align with human preferences beyond mere correctness.

How Machines "Understand": Embeddings and Semantic Space

LLMs don't "understand" language in any human sense, but they create sophisticated mathematical representations that capture semantic relationships. This is achieved through **embeddings**—the transformation of discrete words into continuous high-dimensional vectors.

The Vector Representation

Every word is converted into a list of hundreds or thousands of numbers (a vector). For example, "cat" might become [0.2, -0.5, 0.8, 0.1, ...] in a 768-dimensional space. These numbers encode various semantic properties: meaning, grammatical function, typical context, and more.



In this "semantic space," words with similar meanings occupy nearby positions. The vector distance between "apple" and "fruit" is small, whilst the distance between "apple" and "car" is large. Remarkably, mathematical operations on these vectors can capture semantic relationships: $\text{vector}(\text{"king"}) - \text{vector}(\text{"man"}) + \text{vector}(\text{"woman"}) \approx \text{vector}(\text{"queen"})$.

This geometric representation of meaning—whilst lacking true understanding—enables LLMs to capture nuanced semantic relationships, perform analogical reasoning, and generate contextually appropriate text with surprising sophistication.

Technical Risks: Limitations of the Technology



Hallucinations

LLMs are fundamentally "stochastic parrots"—they prioritise fluent, probable-sounding text over factual accuracy. They can confidently generate convincing falsehoods, fabricate citations, or create entirely fictional events with no indication of uncertainty.



Prompt Injection

Malicious users can "trick" AI systems to bypass safety filters or reveal system prompts through carefully crafted inputs. This vulnerability remains an active research challenge with no complete solution.



Security Vulnerabilities

From data poisoning during training to adversarial attacks at inference time, LLMs present novel security challenges that traditional software defences may not adequately address.



The Fundamental Limitation

LLMs optimise for probability distributions over sequences, not truth. They lack grounding in reality, fact-checking mechanisms, or genuine understanding of whether their outputs correspond to actual states of the world.

Ethical Concerns: Bias, Fairness, and Impact

Beyond technical limitations, LLMs raise profound societal concerns that demand careful consideration from developers, deployers, and policymakers. These systems don't merely reflect existing societal patterns—they can amplify them in ways that perpetuate or even worsen existing inequalities.

Bias Amplification

If training data contains societal biases—and internet data inevitably does—the model learns and reproduces these biases. Gender stereotypes, racial prejudices, and cultural assumptions become embedded in the model's representations and outputs.

Example: A model might associate "doctor" primarily with male pronouns or "nurse" with female pronouns, reflecting historical biases in the training corpus.

Environmental Impact

The carbon footprint of training large models is substantial. Training GPT-4-scale models emits greenhouse gases equivalent to several hundred petrol-powered cars driven for an entire year. As models grow larger and training becomes more frequent, environmental sustainability becomes a pressing ethical concern.

Economic Displacement

As LLMs become capable of performing knowledge work tasks, questions arise about labour market disruption, skill devaluation, and the need for societal adaptation to technological unemployment in cognitive domains.

Misinformation Potential

The ability to generate convincing text at scale creates unprecedented opportunities for automated disinformation campaigns, deepfake content, and manipulation of public discourse.

The Stochastic Parrot: Understanding Fundamental Limitations



The term "stochastic parrot," coined by Bender et al. (2021), captures a crucial insight: LLMs are sophisticated pattern-matching systems that generate text based on statistical regularities in training data—not systems that possess genuine understanding or knowledge.

What This Means in Practice

- **No Truth-Seeking:** Models optimise for plausible-sounding text, not factual accuracy
- **No Comprehension:** They manipulate symbols without grasping underlying meaning or real-world referents
- **No Reasoning:** Apparent reasoning is pattern matching, not logical inference
- **No Intentionality:** Models have no goals, beliefs, or understanding of their own outputs

Understanding LLMs as stochastic parrots—however sophisticated—helps calibrate expectations and informs appropriate use cases whilst avoiding anthropomorphisation that overstates their capabilities.

Essential Takeaways for Students and Researchers

1 Prediction, Not Cognition


LLMs don't genuinely "know" facts or "understand" concepts. They calculate the probability of token sequences based on patterns learned from training data. This distinction is fundamental to understanding both their capabilities and limitations.

2 Context is King

The Transformer's revolutionary ability to link distant words through self-attention mechanisms creates the compelling illusion of human-like thought and comprehension. This architectural innovation, more than raw scale, explains their remarkable performance.


3 Human-in-the-Loop Necessity

Modern LLMs are only as polite, helpful, and appropriate as they are because of thousands of hours humans spent grading their behaviour through SFT and RLHF. Without human feedback, these systems would be knowledgeable but practically unusable.



Future Directions and Open Challenges

Despite remarkable progress, numerous fundamental challenges remain in LLM research and development. Addressing these will shape the next generation of language technologies and their integration into society.



Reasoning and Planning

Improving genuine multi-step reasoning, causal understanding, and long-horizon planning beyond pattern matching




Bias Mitigation

Developing robust techniques to identify, measure, and reduce harmful biases whilst preserving model capability



Efficiency

Creating more parameter-efficient and computationally sustainable training and inference methods



Factual Grounding

Connecting models to external knowledge sources and developing mechanisms for verifiable, attributable outputs



Interpretability

Understanding what models learn, how they make decisions, and why they sometimes fail catastrophically



Value Alignment

Ensuring AI systems reliably pursue goals aligned with human values and societal wellbeing

Conclusion: Responsible Innovation in AI

Large Language Models represent a transformative technology with profound implications for research, industry, and society. Their remarkable capabilities in processing and generating human-like text open extraordinary opportunities across domains from scientific discovery to creative expression.

However, these capabilities come with equally profound responsibilities. As researchers, developers, and informed users, we must maintain clear-eyed understanding of both their potential and their limitations. LLMs are powerful tools—not sentient entities, not infallible oracles, not replacements for human judgement.

Principles for Responsible Use

- Verify critical information independently
- Remain aware of bias and fairness concerns
- Consider environmental and societal impacts
- Maintain human oversight in consequential decisions
- Engage with ongoing ethical discussions

The Path Forward

Progress requires multidisciplinary collaboration between computer scientists, ethicists, policymakers, and domain experts. Technical innovation must proceed hand-in-hand with thoughtful consideration of societal implications.

The future of LLMs depends not only on architectural breakthroughs and scaling laws, but on our collective commitment to developing and deploying these systems responsibly, transparently, and in service of broad human flourishing.